

ABC EVENTS — PROMOTER AI

Scope Document — Discovery Phase Deliverable

1.1 Customer Overview

Company Name	ABC Events
Industry	Live Entertainment — Touring & Concert Promotion
Company Size / Locations	Reasonable Assumption: a multi-market, enterprise-scale North American touring promoter organization with a distinct senior-leadership review layer — sized this way based on two explicit PRD signals, not guessed independently: the Scalability NFR requires "concurrent usage across the full ABC Events promoter base," and the Availability NFR scopes operating hours to "North American time zones" (plural). Exact headcount and venue-partner count are not stated and should be confirmed with ABC Events in the first discovery session — this is a scoping input, not a blocking gap.
Current Tech Stack	An internal, proprietary tour business planning system (system of record for business plans, revenue/expense forecasts, and approvals). Promoter AI is scoped to integrate with this system for live plan context (Section 9, In Scope). No other systems are named in the PRD; vendor should request a data-integration and API capability assessment of this system as an early discovery deliverable, since Section 12 flags this integration as an unconfirmed assumption, not a guaranteed capability.
Primary Stakeholders	Touring Promoters (primary/core users); Senior Leadership Reviewers (secondary users, plan approvers, and expertise-layer authors); ABC Events AI & Data Engineering (document owner, system integrator, expertise-layer curator). Not named in the PRD but required for delivery and flagged here as a gap to close in discovery: Legal/Compliance (outbound communication review, data-residency sign-off), Finance (vendor spend ownership, ROI tracking against Section 13 metrics), and IT/Security (access-control and infrastructure approval for Section 11). Indirect stakeholder: artists' agents and managers, who never use the system but are the recipients of AI-drafted pitch communications and are therefore exposed to its output quality.
Engagement Type	Vendor solution proposal — Discovery Phase. ABC Events is soliciting external vendor proposals for an end-to-end build of the Promoter AI capability set described in the PRD (PRD Section 1).
Document Prepared By	[Candidate Name] — Business Analyst / Solution Consultant
Date	24 July 2026

1.2 Pain Points

Business Pain Points

- **Revision bottleneck at leadership review** — senior leaders spend significant time reviewing tour business plans that require multiple revision cycles, which bottlenecks the approval pipeline and diverts leadership attention from decision-ready work (PRD, Problem Statement).
- **Deal outcomes tied to expertise concentration** — because negotiation expertise sits with a small group of senior leaders, deal quality is uneven across the promoter base rather than being a consistent organizational capability.
- **Competitive exposure from slow plans** — artist routing windows are competitive; a promoter who cannot produce a defensible plan quickly risks losing the deal to a competing promoter or venue (Section 5.1).

Operational Pain Points

- **Escalation-dependent guidance** — promoters routinely escalate to senior leaders for guidance that could, in principle, be answered by reference to historical precedent, consuming senior leadership time that scales poorly across a growing promoter base.
- **Simultaneous multi-deal negotiation load** — promoters negotiate multiple deals concurrently across different artists, markets, and venues, with no system to help prioritize or calibrate effort across them.
- **No structured pre-submission check** — there is no mechanism today for a promoter to self-assess a plan's likely defensibility before it reaches leadership, so weaknesses are caught late, after time has already been invested.

Technology Pain Points

- **No live linkage between planning system and expertise** — the internal tour business planning system captures the plan itself but has no connective layer to historical deal actuals or codified expertise at the point of authoring.
- **Historical deal data is not surfaced as a comparison tool** — historical deal actuals exist but are not organized or exposed in a way that lets a promoter find and compare relevant precedent in the flow of work.

Knowledge Management Pain Points

- **Tacit, undocumented expertise** — feedback on a returned plan is typically qualitative and rooted in a reviewer's personal experience (e.g., "the guarantee is too high for this market") rather than codified, reusable criteria.
- **No structured counterparty knowledge** — negotiation history and behavioral patterns of specific external agents and managers are not captured or shared systematically, so each promoter effectively relearns this per relationship.
- **Knowledge does not compound** — there is currently no repository designed to accumulate and refine deal-making knowledge over time as more deals close and more feedback is generated (PRD Goal 4).

AI Readiness Pain Points

- **Unvalidated data sufficiency** — the PRD assumes historical deal actuals are "sufficiently rich" to ground outputs across artist tiers, markets, and venue types (Section 12); this is an assumption, not a confirmed fact, and is a readiness risk for any AI capability.
- **Dependency on leadership participation** — the codified expertise layer can only be seeded through structured knowledge elicitation sessions with senior leaders, whose time commitment is assumed but not guaranteed (Section 12).
- **Cold-start coverage gaps** — the PRD does not address what happens when a promoter is structuring a deal for a debut artist, a brand-new venue, or a first-time counterparty with no historical comparable data; without an explicit fallback, this is where the "do not fabricate, signal low confidence" requirement will be tested hardest.
- **Emerging-market equity risk** — markets and artist tiers with fewer historical deals will systematically produce lower-confidence, less useful guidance than well-covered markets, which risks the system feeling most helpful exactly where promoters already have the most support — and least helpful where newer promoters and smaller markets need it most.

Reporting Pain Points

- **No visibility into revision-cycle trends** — leadership currently lacks a structured, ongoing view of how many plans are being returned, why, and whether that is improving over time.
- **No traceability reporting** — there is no existing way to confirm that guidance given to a promoter is actually grounded in real data versus generic advice — a requirement the PRD explicitly elevates (Accuracy & Grounding NFR).

User Experience Pain Points

- **Fragmented workflow** — today, calibrating a plan, interpreting a counter-offer, checking counterparty history, and drafting communications are separate, manual, and unsupported activities rather than one connected experience.
- **Uneven experience by tenure** — earlier-career promoters ("strong hustle but limited deal pattern recognition") have materially less support than experienced promoters, with no system narrowing that gap (Section 5.1).

1.3 Specific Requirements / Requests

Functional Requirements

- Find Similar Deals: surface ranked, comparable historical deals (matched on artist tier, market, venue type, and demand profile) so a promoter can calibrate revenue and expense assumptions against actual performance.
- Evaluate Deal Terms: assess proposed guarantee level, venue cost structure, and sellthrough assumptions against codified benchmarks, flagging any element outside a defensible range.
- Interpret Counter-Offer: analyze an incoming response from an artist's agent or manager, surface the underlying position, and suggest constructive next moves.
- Counterparty History: return prior deal history and negotiation patterns for a specific external agent or manager to inform negotiation approach.
- Draft Pitch Communication: generate an outbound pitch communication grounded in the current plan context and relevant historical precedent.
- Frame Internal Ask: help a promoter structure a request to leadership, informed by what leadership has historically looked for (and pushed back on) in plan approvals.
- Comprehensive Plan Review: compose the above capabilities into a single holistic pre-submission assessment of a business plan.
- Live plan context retrieval: read the promoter's active business plan from the internal tour business planning system, scoped to the authenticated promoter, to ground every capability response.
- Cold-start handling: when a deal involves a debut artist, a new venue, or a first-time counterparty with insufficient historical data, every relevant capability must explicitly state that precedent is limited, show the closest available adjacent-tier comparisons (clearly labeled as lower-confidence substitutes), and route the promoter toward senior-leader input rather than presenting a confident but thinly-supported answer.
- Human-review gate for outbound content: any AI-drafted pitch communication to an artist's representatives must be presented as an editable draft requiring explicit promoter (and, where policy requires, Legal) review before it is sent; the system must never transmit outbound communications autonomously.
- Plan revision traceability: every submission, review decision, and return reason for a business plan must be recorded as a discrete, queryable revision record — not just a running counter — since revision-cycle reduction is a core success metric (Section 13) and must be auditable at the individual-plan level.

Non-Functional Requirements

- Performance: capability responses returned within 5 seconds for a standard single-capability query under normal load.
- Availability: system available during ABC Events business hours across North American time zones, with a target uptime of $\geq 99.5\%$.
- Scalability: must support concurrent usage across the full ABC Events promoter base without performance degradation.
- Accuracy & Grounding: every recommendation must be traceable to a historical comparable deal or a codified expertise principle; the system must not fabricate a response when data is insufficient and must explicitly signal low confidence instead.
- Latency (plan context): retrieving a promoter's active plan for context must not materially add to total response time.
- Expertise layer maintainability: the codified expertise layer must be updatable by ABC Events without a full system redeployment.
- Role-based access control: distinct permission tiers for Promoter, Leadership Reviewer, and Expertise Curator roles, since "promoters may only access data scoped to their own plans" (Section 11) implies role differentiation the PRD names but does not structurally define.

- Disaster recovery: a defined Recovery Time Objective and Recovery Point Objective for both the plan-context integration and the historical-deal data store, consistent with the $\geq 99.5\%$ uptime target — the PRD sets an availability bar but not a recovery standard if that bar is breached.
- Fairness monitoring: recommendation confidence and grounding rates must be tracked and reported by market and artist tier, not only in aggregate, so that systematically under-served markets (thin historical data) are visible to leadership rather than hidden inside a global average.

Delivery Phasing (Reasonable Assumption)

The PRD states the capability set is "to be built progressively" (Section 7) but does not specify sequence. The phasing below sequences by data-maturity and output-risk — lower-risk, better-grounded capabilities first; capabilities that generate outbound content or depend on richer counterparty data last.

- **Phase 1 — Foundation:** Find Similar Deals, Evaluate Deal Terms, and the core historical-deal / live-plan data integration. Highest leverage on the revision-cycle and review-time success metrics, and the lowest workflow risk (read-only, advisory).
- **Phase 2 — Structured self-review:** Comprehensive Plan Review, backed by Plan_Revision_History reporting, once Phase 1 grounding data is validated in production.
- **Phase 3 — Counterparty intelligence:** Interpret Counter-Offer and Counterparty History, gated on the Counterparty entity holding sufficient deal history to avoid cold-start guidance on live negotiations.
- **Phase 4 — Outbound communication:** Draft Pitch Communication and Frame Internal Ask, sequenced last because these capabilities generate content that leaves the system and carries the highest reputational and legal exposure if ungrounded.

Integration Requirements

- Bi-directional or read integration with ABC Events' internal tour business planning system to retrieve live, promoter-scoped plan context.
- Ingestion pipeline from historical deal actuals (revenue forecast vs. actual, expense structure vs. actual, sellthrough vs. actual) into the comparable-deals data layer.
- A structured intake path for senior-leader knowledge elicitation output to feed the codified expertise layer, with human curation before anything is published for use.

Security Requirements

- Access control: promoters may only access data scoped to their own plans and authorized deal history.
- Data residency: no sensitive deal, artist, or plan data may leave ABC Events' managed infrastructure.
- Full interaction logging: all inputs, outputs, and feedback signals logged with sufficient detail for post-hoc audit and debugging.

Compliance Requirements

- Knowledge handling: expertise layer contributions must be reviewed and curated before use; no unreviewed material may be incorporated into a live capability response.
- Data scope exclusion: no customer or fan data may be included in any capability context (touring-deal data only).

Future Enhancements

- PRD Goal 5 explicitly frames this build as "a foundation for additional promoter-facing AI capabilities" — future phases may extend beyond deal structuring into adjacent promoter workflows (Reasonable Assumption: specific future capabilities are not named in the PRD).

Out of Scope

- Fully automated plan approval without leadership review — all capability outputs remain advisory; leadership retains final approval authority (Section 9 / Section 12).

- Live negotiation with external parties — the system coaches the promoter; it does not interact with agents or managers directly.
- Replacement of existing scheduling or expense tooling.
- Multi-language support.

Dependencies

- Sufficiency and quality of historical deal actuals across a meaningful range of artist tiers, markets, and venue types (Section 12, Assumption).
- The internal tour business planning system's ability to expose live, promoter-scoped plan data to an external capability layer (Section 12, Assumption).
- Senior leadership's committed time for structured knowledge elicitation sessions to seed the codified expertise layer (Section 12, Assumption).
- Early-adopter promoter cohort willing to give candid feedback during validation, since the feedback loop is how the expertise layer improves post-launch (Section 12, Assumption).

1.4 Success Criteria

- **Fewer plans need major revision** — +25% increase (vs. baseline) in plans reaching leadership review with no further major revision required, measured within 6 months of launch, directly reducing rework for both promoters and leadership.
- **Fewer revision cycles per plan** — -30% reduction in average revision cycles per plan (vs. baseline) within 6 months, indicating promoters are submitting better-calibrated plans on the first pass.
- **Faster leadership review** — -40% reduction in leadership review time per plan (vs. baseline) within 6 months, freeing senior leadership capacity for decision-ready work rather than qualitative coaching.
- **Sustained promoter trust** — promoter trust score ("would use again") of ≥80%, sustained — evidence that guidance is genuinely useful rather than adopted out of obligation.
- **Full recommendation grounding** — 100% of recommendations traceable to a historical comparable deal or a codified expertise principle, with explicit low-confidence signaling when data is insufficient — the non-negotiable trust and auditability bar for the system.
- **Leadership confidence in the reasoning, not just the outcome** — ≥80% sustained leadership agreement with capability reasoning, confirming the system's logic — not only its recommendations — holds up under senior scrutiny.

1.5 Database Design

Logical data model. Each row states purpose, keys, columns, relationships, validation rules, ownership, and retention, sized to support the capability set defined in PRD Section 7.

Category	Database Name	Column Definitions
Database	Markets	Purpose: Normalizes market/region references shared by Artists, Venues, and Historical_Deals so tier benchmarks compare like-for-like instead of relying on free-text market names. Primary Key: market_id Foreign Keys: None (top-level reference table) Columns: market_id, market_name, region, timezone, currency_code, market_size_tier Relationships: One-to-many with Artists, Venues, Historical_Deals

Category	Database Name	Column Definitions
		Validation Rules: currency_code must be a valid ISO 4217 code; timezone constrained to the North American set covered by the Availability NFR Owner: ABC Events AI & Data Engineering Retention: Retained indefinitely (reference data)
Database	Artists	Purpose: Master record for every artist ABC Events tours or evaluates; anchors artist-side context for deals, plans, and benchmarks. Primary Key: artist_id Foreign Keys: home_market_id → Markets; primary_agent_id → Counterparties (nullable) Columns: artist_id, artist_name, genre, current_artist_tier, home_market_id, primary_agent_id, debut_artist_flag, first_toured_date, created_at, updated_at Relationships: One-to-many with Historical_Deals, Business_Plans; many-to-one with Markets Validation Rules: current_artist_tier constrained to the defined tier enum; changes to current_artist_tier require Senior Leadership sign-off, since tier drift affects live benchmarking Owner: ABC Events AI & Data Engineering; artist_tier changes approved by Senior Leadership Retention: Retained indefinitely (master data). Note: Historical_Deals stores its own deal-time tier snapshot rather than joining live to this table, so a later tier change never silently rewrites past deal comparisons.
Database	Venues	Purpose: Master record of venues, their capacity, tier, and cost-structure characteristics — the active economic variable described in PRD Section 2.4. Primary Key: venue_id Foreign Keys: market_id → Markets Columns: venue_id, venue_name, venue_tier, capacity, market_id, base_rent_structure, production_compatibility_notes, ancillary_revenue_share_pct, opened_date, is_active Relationships: One-to-many with Historical_Deals, Business_Plans; many-to-one with Markets Validation Rules: capacity > 0; venue_tier constrained to {theater, amphitheater, arena, stadium} Owner: ABC Events AI & Data Engineering Retention: Retained indefinitely; decommissioned venues are marked is_active = false rather than deleted, to preserve historical comparable integrity
Database	Promoters	Purpose: Master record of touring promoters; anchors access-control scoping and the tenure-based segmentation used in adoption analytics. Primary Key: promoter_id Foreign Keys: home_market_id → Markets Columns: promoter_id, promoter_name, tenure_start_date, experience_level, home_market_id, active_flag

Category	Database Name	Column Definitions
		Relationships: One-to-many with Business_Plans, Historical_Deals, Feedback_Signals Validation Rules: tenure_start_date cannot be in the future; experience_level is system-derived from tenure and deal count, not self-declared, to keep segmentation objective Owner: ABC Events AI & Data Engineering; row-level access enforced per Section 11 (promoters see only their own scoped data) Retention: Retained while employed; anonymized on departure, retained 1 additional year for trend continuity
Database	Senior_Leaders	Purpose: Identifies the individuals whose tacit expertise seeds and approves the Codified Expertise Layer, and who review submitted plans. Primary Key: leader_id Foreign Keys: None Columns: leader_id, leader_name, role_title, market_expertise_tags, elicitation_availability_status Relationships: One-to-many with Knowledge_Elicitation_Sessions and Plan_Revision_History (as reviewer); referenced by Codified_Expertise_Layer as source and reviewer Validation Rules: elicitation_availability_status must be kept current — Section 12 names leader time commitment as an assumption, and this field is how that risk is tracked operationally rather than assumed away Owner: ABC Events AI & Data Engineering; self-maintained by leadership Retention: Retained while active in role
Database	Knowledge_Elicitation_Sessions	Purpose: Logs each structured session used to seed or update the expertise layer, giving every principle a traceable origin (Section 8). Primary Key: session_id Foreign Keys: leader_id → Senior_Leaders Columns: session_id, leader_id, session_date, topic_area, facilitator, status, resulting_entry_ids Relationships: One-to-many with Codified_Expertise_Layer Validation Rules: status constrained to {scheduled, completed, cancelled}; resulting_entry_ids may only reference entries with review_status ≥ curated Owner: ABC Events AI & Data Engineering Retention: Retained indefinitely — this is the provenance record for every expertise-layer entry
Database	Counterparties	Purpose: Negotiation profiles for external agents/managers; powers Counterparty History. Primary Key: agent_id Foreign Keys: Referenced by Agent_Artist_Representation and Historical_Deals Columns: agent_id, agent_name, agency_name, typical_guarantee_stance, historical_deal_count, avg_revision_cycles, negotiation_pattern_notes, deal_history_sufficiency_flag, last_interaction_date

Category	Database Name	Column Definitions
		Relationships: One-to-many with Historical_Deals and Agent_Artist_Representation Validation Rules: deal_history_sufficiency_flag is automatically set to false when historical_deal_count falls below a defined minimum — this is the field the Negotiation Support Agent checks before treating a counterparty profile as reliable; negotiation_pattern_notes may only be populated from curated Codified_Expertise_Layer entries, never freeform Owner: ABC Events AI & Data Engineering Retention: Retained while active in market; archived after 3 years of no recorded activity
Database	Agent_Artist_Representation	Purpose: Junction table normalizing the many-to-many relationship between agents/managers and the artists they represent — representation changes over time and one agent commonly represents multiple artists, which a single list field on Counterparties could not model correctly. Primary Key: representation_id Foreign Keys: agent_id → Counterparties; artist_id → Artists Columns: representation_id, agent_id, artist_id, start_date, end_date, is_current Relationships: Many-to-one with Counterparties and Artists Validation Rules: Only one is_current = true row permitted per artist_id at any time Owner: ABC Events AI & Data Engineering Retention: Retained indefinitely — representation history directly informs negotiation-pattern interpretation
Database	Historical_Deals	Purpose: Stores closed touring deal records with forecast-vs-actual performance; core comparable-deal corpus for Find Similar Deals and Evaluate Deal Terms. Primary Key: deal_id Foreign Keys: artist_id → Artists; venue_id → Venues; market_id → Markets; promoter_id → Promoters; agent_id → Counterparties Columns: deal_id, artist_id, venue_id, market_id, promoter_id, agent_id, event_date, artist_tier_at_deal_time, venue_tier_at_deal_time, guarantee_amount, projected_sellthrough_pct, actual_sellthrough_pct, projected_revenue, actual_revenue, ancillary_revenue_actual, vip_package_revenue_actual, breakeven_point, deal_outcome, deal_cancelled_flag, revision_cycle_count, created_at, updated_at Relationships: Many-to-one with Artists, Venues, Markets, Promoters, Counterparties; one-to-many with Deal_Expense_Line_Items Validation Rules: $0\% \leq \text{sellthrough_pct} \leq 100\%$; guarantee_amount > 0; artist_tier_at_deal_time and venue_tier_at_deal_time are snapshots captured at deal close, never re-derived from current master-data tier; deal_cancelled_flag = true records are excluded from comparable-deal matching by default (a deal approved but never performed should not silently inflate the precedent corpus) Owner: ABC Events AI & Data Engineering

Category	Database Name	Column Definitions
		Retention: Retained indefinitely as the core comparable corpus; relevance reviewed annually
Database	Deal_Expense_Line_Items	Purpose: Normalized expense detail per deal (artist guarantee, venue costs, production, marketing, ticketing/operations), replacing a single unstructured expense field so category-level benchmarking and analytics are actually queryable. Primary Key: line_item_id Foreign Keys: deal_id → Historical_Deals Columns: line_item_id, deal_id, expense_category, projected_amount, actual_amount Relationships: Many-to-one with Historical_Deals Validation Rules: expense_category constrained to {artist_guarantee, venue_costs, production, marketing, ticketing_ops}; amounts non-negative Owner: ABC Events AI & Data Engineering Retention: Same retention as parent Historical_Deals record
Database	Business_Plans	Purpose: Read-scoped mirror of live plan context from the internal tour business planning system; grounds in-flight capability responses (Section 9). Primary Key: plan_id Foreign Keys: promoter_id → Promoters; artist_id → Artists; venue_id → Venues Columns: plan_id, promoter_id, artist_id, venue_id, plan_status, proposed_guarantee, proposed_sellthrough_pct, revenue_projection_total, expense_projection_total, venue_cost_structure, submission_date, last_synced_at, sync_status, leadership_reviewer_id Relationships: One-to-many with Feedback_Signals and Plan_Revision_History; synced (not owned) from the internal planning system Validation Rules: plan_status constrained to {draft, submitted, approved, returned}; sync_status constrained to {synced, stale, sync_error} — capability responses must flag or suppress guidance whenever sync_status ≠ synced, so promoters are never advised against out-of-date plan data Owner: ABC Events AI & Data Engineering (source of truth remains the internal planning system) Retention: Plan lifecycle plus 3 years for audit purposes
Database	Plan_Revision_History	Purpose: Captures every individual submission and review decision for a plan as a discrete record. This is the entity the PRD's core success metrics — revision cycles per plan, leadership review time — are actually measured against, so it cannot be a running counter alone. Primary Key: revision_id Foreign Keys: plan_id → Business_Plans; reviewer_id → Senior_Leaders (nullable until reviewed) Columns: revision_id, plan_id, revision_number, submitted_at, reviewed_at, reviewer_id, decision, returned_reason_category, fields_changed

Category	Database Name	Column Definitions
		Relationships: Many-to-one with Business_Plans and Senior_Leaders Validation Rules: revision_number strictly increasing per plan_id; decision constrained to {approved, returned, pending}; returned_reason_category required whenever decision = returned, so return reasons are analyzable rather than free text Owner: ABC Events AI & Data Engineering Retention: Matches parent Business_Plans record: plan lifecycle plus 3 years
Database	Codified_Expertise_Layer	Purpose: Structured principles, benchmarks, counterparty patterns, and leadership approval criteria that ground every capability output (Section 8). Primary Key: expertise_entry_id Foreign Keys: source_leader_id → Senior_Leaders; reviewed_by_leader_id → Senior_Leaders; elicitation_session_id → Knowledge_Elicitation_Sessions Columns: expertise_entry_id, entry_type, artist_tier, venue_tier, market_id, acceptable_range_min, acceptable_range_max, red_line_threshold, source_leader_id, reviewed_by_leader_id, elicitation_session_id, review_status, version, effective_date, superseded_by Relationships: Many-to-one with Senior_Leaders (as both author and reviewer) and Knowledge_Elicitation_Sessions Validation Rules: Only entries with review_status = curated or published may be served to a capability at runtime; reviewed_by_leader_id must differ from source_leader_id for any entry that changes an existing red-line threshold, enforcing segregation of duties on the most consequential edits Owner: ABC Events AI & Data Engineering, curated jointly with Senior Leadership Retention: Versioned; superseded entries archived, never deleted, to preserve the audit trail
Database	Feedback_Signals	Purpose: Captures promoter and leadership ratings, overrides, and annotations on capability outputs — the mechanism that refines the expertise layer over time (Section 8 / Goal 4). Primary Key: feedback_id Foreign Keys: plan_id → Business_Plans; expertise_entry_id → Codified_Expertise_Layer (nullable) Columns: feedback_id, capability_name, plan_id, expertise_entry_id, submitted_by, role, rating, override_flag, override_reason, capability_confidence_at_time, annotation_text, submitted_at, incorporated_flag Relationships: Many-to-one with Business_Plans and Codified_Expertise_Layer Validation Rules: rating between 1 and 5; override_reason required whenever override_flag = true; capability_confidence_at_time captured so leadership can later analyze whether overrides cluster around low-confidence outputs (expected and healthy) or high-confidence ones (a quality signal worth investigating) Owner: ABC Events AI & Data Engineering

Category	Database Name	Column Definitions
		Retention: 5 years, to support longitudinal auditing of expertise-layer evolution
Database	Interaction_Audit_Log	Purpose: Full log of every capability interaction, satisfying the Auditability NFR (Section 10 / 11). Primary Key: interaction_id Foreign Keys: promoter_id → Promoters; plan_id → Business_Plans Columns: interaction_id, promoter_id, plan_id, capability_name, input_payload, output_payload, confidence_score, grounding_reference_ids, timestamp, response_time_ms Relationships: Many-to-one with Promoters and Business_Plans Validation Rules: response_time_ms flagged automatically if it exceeds the 5-second performance NFR Owner: ABC Events AI & Data Engineering Retention: Minimum 2 years, per audit requirement

1.6 Analytics Design

Each insight is framed to enable a specific promoter or leadership decision, not as a vanity metric.

Category	Insight Name	Visualization Type	Logic / Description
Analytics	Plan Revision Rate Trend	Line Chart	Tracks % of plans requiring at least one major revision over time, plotted against a rolling baseline. Filters: market, artist tier, promoter tenure. Drilldown to individual plan history. Enables leadership to confirm progress against the +25%/-30% success targets and spot regressions early.
Analytics	Guarantee vs. Sellthrough Defensibility	Quadrant / Scatter Chart	Plots each active plan's proposed guarantee against its projected sellthrough, overlaid with the acceptable-range band from the Codified Expertise Layer for that artist/venue tier. Points outside the band are flagged. Drilldown to the specific benchmark principle driving the flag. Lets promoters and leadership see, at a glance, which in-flight plans are at revision risk before submission.
Analytics	Comparable Deal Performance	Table with Drilldown	Ranked list of historical deals matched to the current plan's artist tier, venue tier, and market, showing projected-vs-actual sellthrough and revenue variance. Filters: market, date range, artist tier. Drilldown to full deal record. Directly supports promoters calibrating revenue and expense assumptions against real precedent.
Analytics	Leadership Review Time by Segment	Bar Chart	Average leadership review time per plan, segmented by reviewer, artist tier, and market. Filters: time period, reviewer. Drilldown to individual plan review duration. Directly measures progress toward the

Category	Insight Name	Visualization Type	Logic / Description
			–40% review-time success target and surfaces where review effort concentrates.
Analytics	Recommendation Grounding & Confidence	Card / Gauge	Displays the percentage of capability outputs that are fully traceable to a historical comparable or codified principle (target: 100%), alongside the rate of explicit low-confidence signaling. Filters: capability name, date range. Gives leadership a direct, ongoing trust and compliance signal for the system itself.
Analytics	Counterparty Negotiation Pattern Summary	Table / Heatmap	Summarizes each agent/manager's historical guarantee stance, average revision cycles, and deal-close rate. Filters: agency, artist tier. Drilldown to underlying Historical_Deals records. Informs a promoter's negotiation strategy before engaging a specific counterparty.
Analytics	Promoter Trust & Adoption	Card + Trend Line	Tracks the sustained promoter trust score ("would use again") against the ≥80% target, segmented by promoter tenure and market using the Promoters entity's experience_level field. Filters: cohort, time period. Signals to leadership where additional training or expertise-layer refinement may be needed.
Analytics	Expertise Layer Freshness	Bar Chart / Table	Shows the proportion of Codified Expertise Layer entries updated within the last quarter versus entries flagged stale, segmented by entry type (principle, benchmark, counterparty pattern, approval criterion). Drilldown to individual entries pending review. Ensures the knowledge base is compounding, not decaying, per Goal 4.
Analytics	Data Coverage / Cold-Start Gap Map	Heatmap	Cross-tabulates artist tier × venue tier × market, showing the count of historical comparable deals available in each cell, with cells below the minimum-sufficiency threshold visibly flagged. Drilldown to which specific artist/venue combinations are driving a gap. Directly operationalizes the Accuracy & Grounding NFR and the emerging-market equity risk identified in Pain Points — tells leadership exactly where the system will (and won't) be able to ground a confident answer, before a promoter finds out the hard way mid-negotiation.
Analytics	Feedback Override & Disagreement Rate	Bar Chart / Table	Rate at which promoters or leadership override a capability's recommendation, segmented by capability and by the confidence_at_time of the original output. Filters: capability, role, date range. Distinguishes healthy disagreement (overrides concentrated in already-low-confidence outputs) from a quality problem (overrides concentrated in

Category	Insight Name	Visualization Type	Logic / Description
			high-confidence outputs) — the single clearest leading indicator that a specific expertise-layer entry needs review before it shows up in a satisfaction-score decline.

1.7 Agent Design

Agents are grouped around business workflows, not 1:1 with individual PRD capabilities, so each has a coherent trigger and a single clear output. All agents remain advisory — none bypasses leadership approval authority (PRD Section 9, Out of Scope).

Category	Agent Name	Workflow (Trigger → Output)	CTAs → Agent Workflow
Agent	Deal Calibration Agent	Trigger: promoter is actively editing the revenue or expense section of a business plan in the planning system. → Workflow: retrieves live plan context, queries Historical_Deals (excluding cancelled deals) for matches on artist tier/venue tier/market, evaluates proposed guarantee and sellthrough against Codified_Expertise_Layer benchmarks. Cold-start case: if fewer than the defined minimum number of comparable deals exist for the exact tier/market combination, the agent widens to adjacent tiers, labels those results as "lower-confidence, adjacent-tier substitutes," and does not present a single-point recommendation. → Output: ranked comparable deals plus a defensibility assessment with an explicit confidence score.	"Show me comparable deals" → Deal Calibration Agent (Find Similar Deals) "Is my guarantee defensible?" → Deal Calibration Agent (Evaluate Deal Terms) "Explain this benchmark" → drills into the source expertise principle
Agent	Negotiation Support Agent	Trigger: promoter logs or pastes an incoming counter-offer, or requests a counterparty profile. → Workflow: parses the counter-offer against deal context and the Counterparties/Agent_Artist_Representation records; retrieves that agent/manager's negotiation pattern from prior deals. Cold-start case: if Counterparties.deal_history_sufficiency_flag is false (new or rarely-encountered counterparty), the agent states plainly that no reliable pattern exists yet rather than inferring one from a handful of deals. → Output: interpreted counterparty position, historical pattern summary (or an explicit no-pattern notice), and suggested constructive next moves; escalates to the Communication Drafting Agent only if the promoter requests a written response.	"What is this counter-offer really asking?" → Negotiation Support Agent "What's this agent's history?" → Negotiation Support Agent "Help me respond" → escalates to Communication Drafting Agent

Category	Agent Name	Workflow (Trigger → Output)	CTAs → Agent Workflow
Agent	Communication Drafting Agent	Trigger: promoter requests a pitch draft to an artist's representatives, or help framing an internal ask to leadership. → Workflow: pulls current plan context plus relevant historical precedent (for pitches) or Leadership Approval Criteria (for internal asks) and generates an editable draft. → Output: a draft communication with source references the promoter can inspect. Mandatory gate: outbound drafts are never transmitted by the agent — they are held in a draft state requiring explicit promoter edit-and-approve action (and Legal review where policy requires) before being sent through existing communication channels; the agent makes no commitments to the artist's representatives on ABC Events' behalf.	"Draft a pitch to [artist] reps" → Communication Drafting Agent (draft only — human sends) "Help me frame this ask to leadership" → Communication Drafting Agent
Agent	Plan Review Agent	Trigger: promoter requests a full pre-submission review, or the plan status is changed to "ready for submission." → Workflow: composes Deal Calibration output with Leadership Approval Criteria and the plan's own Plan_Revision_History (has this plan already been returned once, and for what reason); assigns a confidence score to each flagged weakness. → Output: a prioritized list of issues to fix before submission, explicitly distinguishing "benchmark violation" (high confidence) from "limited precedent, use judgment" (cold-start) flags; low-confidence or high-stakes items are routed to human escalation rather than resolved automatically, consistent with the "no fully automated approval" scope boundary.	"Review my plan before I submit" → Plan Review Agent "What will leadership push back on?" → Plan Review Agent "Talk to a senior leader" → escalates to human review queue
Agent	Expertise Signal Agent	Trigger: accumulated Feedback_Signals for a given expertise entry or capability cross a defined disagreement or override threshold. → Workflow: aggregates override rates, ratings, and annotation themes into a statistical summary — it does not draft new principles, benchmarks, or wording of any kind. → Output: a prioritized list of "candidate topics for review" (e.g., "guarantee benchmark for mid-tier arena venues in the Southeast market has been overridden 40% of the time in the last quarter") routed to Senior Leadership as input to their next structured knowledge elicitation session. This deliberately stops short of authoring expertise-layer content,	"What's being overridden most?" → Expertise Signal Agent "Flag this benchmark as outdated" → adds a manual signal to the same queue All outputs route to Senior Leadership for the actual elicitation session — never auto-published

Category	Agent Name	Workflow (Trigger → Output)	CTAs → Agent Workflow
		since Section 8 defines that layer as elicited from senior leaders in structured sessions, and Section 11 requires that no unreviewed material — including AI-drafted material — ever be incorporated.	